

FISH

Emil Tullstedt

January 11, 2017

Contents

1	Task Specification	2
2	Development Environment	2
2.1	The implementor's hardware & software	2
2.2	Compiling & Running	2
3	Description of Technical Implemenation	3
3.1	Overview of Classes	3
3.1.1	The story of how classes relate to each other	3
3.2	UDP Multicast for File Discovery	5
3.3	Network Packages	5
3.3.1	Requesting a file listing	5
3.3.2	Informing about the existance of a file	5
3.3.3	Requesting a file download	5
3.4	User Interface	6
4	Running the Application	7
5	Documentation	7

1 Task Specification

FISH is a *distributed* file sharing system for local networks using UDP multicast to search for files shared by a FISH-peer in the network.¹ A file is shared either individually or by being part of the shared directory for the FISH client and is assumed to be uniquely described by filename within the network.

When a file is downloaded from another peer in the network, it will be stored in a specified download destination directory (by default `/tmp`). If the specified download destination directory is equivalent (either by being the same as or a link to) to the shared directory a file downloaded through FISH will cause an automatic renewal of files shared from the specific host.

2 Development Environment

2.1 The implementor's hardware & software

Computer Lenovo ThinkPad X250 with Intel Core i3-5010U @ 2.10GHz and 8 GB DDR3 RAM

Operating System Fedora 23 running Intel 64 Linux 4.8.13

Java JRE & JDK 64-bit OpenJDK 1.8.0 (Java SE 8)

Build system Gradle 2.5

2.2 Compiling & Running

1. Go to the root directory in the *git repository*
2. Make sure the **multicast** repository is the currently checked out repository (by running `git checkout multicast` if necessary)
3. Run `$ gradle build` in order to build the application.
4. Go to *build/classes/main*

¹There is also an implementation where the file searching is done using a centralized server and persistent TCP connections between that server and the peers in the network which is not the subject of this report

5. Run `$ java is.mjuk.fish.Client <file or directory to share>`

The application is now running and should be able to parse the commands **find**, **destination** and **download** to find a file, get (without any arguments) or set the download destination (default `/tmp`) and download a file to the current download destination.

3 Description of Technical Implementation

3.1 Overview of Classes

The entire FISH-implementation is packaged in the *is.mjuk.fish* Java package with separation of concerns on a class-level. The source files are available in `<gitroot>/src/main/java/is/mjuk/fish`.

Client Command-line interface and launcher for the FISH protocol implementation

DatagramHandler Handles multicasting between peers

Downloader Thread which download requested files in the background

Helpers Static functions and constants that can be used for convenience

PeerListener Open a ServerSocket that sends file per request

PeerListener.SinglePeerListener Handles a single client for a PeerListener

UnicastListener Listens to the peer sharing for information about searches

There is also a single interface **ConnectorInterface** which is currently used only by **DatagramHandler**. The **ConnectorInterface** defines a bare minimum for a class running on one or more threads which communicates over one or more sockets.

3.1.1 The story of how classes relate to each other

The application execution has a single entry point from which each and every peer much originate, and that is the **Client**-class. The **Client** starts by creating it's **DatagramHandler** and **Downloader** (in separate threads) for connecting to the peer swarm and preparing for future downloads. They are

both stored as class variables (two each, one for the Thread-object and one for the class itself).

Both these classes gets access to the **Client**'s **this** object in order for the **Downloader** know where to store files and update the file listing after a download and for the **DatagramHandler** object to be able to determine if a file is available in the client's shared file list or not.

The **PeerListener** is started from the **Client** to prepare the peer to send files to requesting peers and handles the TCP connections between the client and other clients. The **PeerListener** is passed with the **this** object for the **Client** in order to be able to fetch a new list of files when necessary. **PeerListener** opens up a random TCP-port between 28000 and 28800 for it's **ServerSocket**.

For every incoming requesting the **PeerListener** defers it's connection socket to a **SinglePeerListener**-thread which is the "active" class within the **PeerListener**-class and handles the actual communication with the other peer. For most intents and purposes, the **PeerListener** name is used for both these two classes interchangeably (as they are simply an implementational detail).

The **UnicastListener** is started and has access to the **Downloader** and receives the port number on which the **PeerListener** listens in order to assign the same port for UDP.

The **UnicastListener** transmits information about file locations it gets access to to the **Downloader** which determines if the data is queued for download or not. If the **Downloader** determines that the file should be downloaded, it starts a download request to a foreign **PeerListener** (or compatible TCP-socket in another implementation of this FISH-protocol).

The sockets opened are in summary:

- An UDP multicast listening socket in **DatagramHandler**
- An UDP unicast sending socket in **DatagramHandler**
- A TCP server socket in **PeerListener**
- An UDP unicast listening socket in **UnicastListener**
- A TCP client socket in **Downloader**

3.2 UDP Multicast for File Discovery

A peer in the network may ask for files over the UDP Multicast channel *239.10.10.10*. If a peer (including the original peer) finds the specified file amongst their shared files, they will respond to the original query using unicast.

The peer who queried for the file may proceed to ignore the information, display it to the user or open a TCP-connection to a selected peer in order to download the file to it's own destination directory. The implementation assumes the user never wants to ask for files for the purpose of ignoring the replies, although that might be abused by a malicious client which is purposefully trying to incite a DDoS of a network if FISH were widely deployed within a network or in a more legitimate use-case used by a client which has successfully downloaded a file and isn't interested by more file listings.

3.3 Network Packages

All communication is serialized over the network as UTF-8 byte packages except for sending the files in themselves (which are transmitted as-is in 4096 byte chunks)

3.3.1 Requesting a file listing

A request for a listing of available sites where a file may be found is sent over the UDP Multicast channel in the format `FIND <unicast port> <file name (may contain spaces)>`.

3.3.2 Informing about the existence of a file

A response to a file listing request takes the form `<TCP Peer port> <file name (may contain space)>` and is sent to the UDP unicast port given in a file listing request.

3.3.3 Requesting a file download

The requester opens a connection to the requestees TCP peer port followed by `DOWNLOAD <file name (may contain spaces)>`. The requestee

responds by sending a 5 bytes status message as described below

E.DNF Did not find the file (and won't send)

K.SEN Found the file. Proceeding to send it to requester

If an error occurred (status message beginning with **E** in case of future expansion of the FISH protocol) the connection is terminated and the requester is assumed to proceed to the next possessor of the file (if possible) ² and not contact the requestee anymore for the same file request. *Note:* The requester should not assume repeated failure if the user initializes a new request and should thus ask the requestee again if so happens.

If the file is successfully found and transmission begins (as indicated by **K.SEN**) the file is sent over the same established TCP socket in incremental chunks of 4096 bytes (statically assigned value) until done. This operation is assumed successful if the connection is gracefully closed.

3.4 User Interface

The user communicates with FISH using the FISH Client which is a simple command-line interface which accepts **find**, **destination** and **download** as described in section 2.2.

destination Display current download destination

destination <path> Set a new download destination

download <file> Send a file listing request for the specified file and download it automatically

exit Close all connections gracefully and die

find <file> Send a file listing request for the specified file

Whenever the client receives the response to a file listing request it displays the filename, hostname and TCP port to the user (whom may then proceed to manually download the file using netcat or a similar tool, by simply stripping the first five bytes of the resulting file).

²I.e. the file is available from more peers. The requester should request the file from the next available peer until successful and this is how my implementation is handling this

4 Running the Application

See section 2.2 "Compiling & Running" and section 3.4 "User Interface" for instructions on how to compile and run the application.

5 Documentation

The implementation's documentation is available in Javadoc-form on

<https://sakjur.mjuk.is/id2212>